

# Profiling and Timing your code

Thomas Hauser

[Thomas.hauser@colorado.edu](mailto:Thomas.hauser@colorado.edu)

Link to survey on this topic: <http://tinyurl.com/rcpresurvey>

[https://github.com/ResearchComputing/Parallelization\\_Workshop/tree/master/Day1](https://github.com/ResearchComputing/Parallelization_Workshop/tree/master/Day1)

# Outline

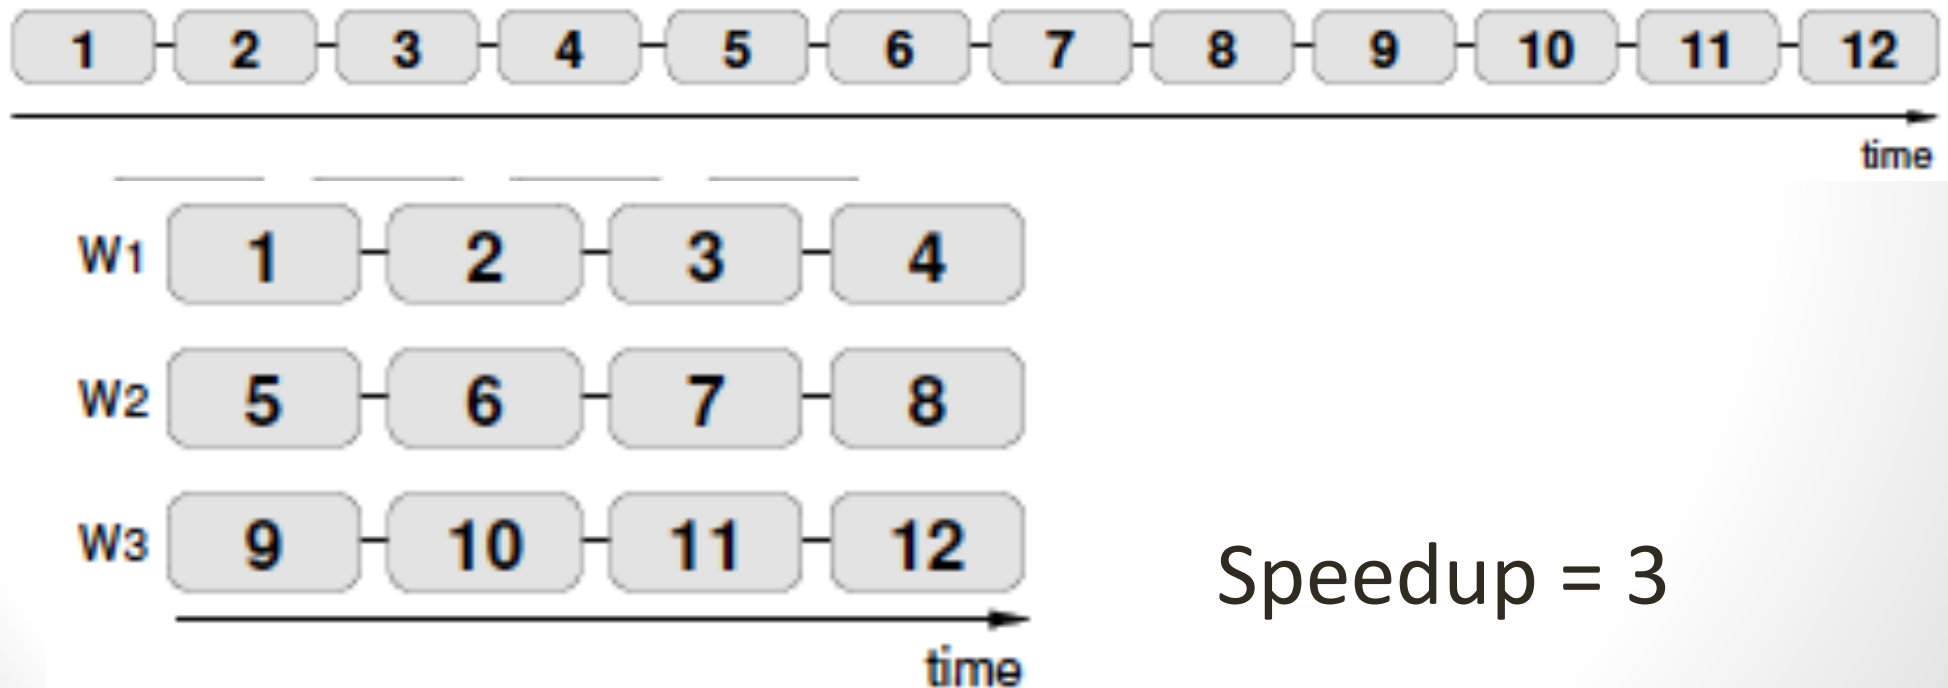
- Why measuring your parallel performance?
- Strong Scaling
- Weak Scaling
- Example measuring your performance
- Memory usage

# Why Measure Performance

- Processors are not getting faster
  - Only performance improvement through parallelism
- Guidance on how many cores to use
- Identify bottlenecks in your code
- Have changes improved the performance of the code
- Speedup
  - Increase of performance processing the same problem
  - Parallel computing – using more cores or nodes

# Speedup Formula

$$\text{Speedup} = \frac{\text{Sequential execution time}}{\text{Parallel execution time}}$$



Speedup = 3


# Execution Time Components

- Inherently sequential computations:  $s(n)$
- Potentially parallel computations:  $p(n)$
- Communication operations:  $c(n, p)$
- Speedup expression:

$$S \leq \frac{s+p}{s(n)+p/N+c}$$

# Parallel Overhead

- Overhead because of
  - Startup time
  - Synchronizations
  - Communication
  - Overhead by libraries, compilers
  - Termination time
- Other barriers to perfect speedup
  - Not perfectly load balanced




# Efficiency

$$\text{Efficiency} = \frac{\text{Sequential execution time}}{\text{Processors} \times \text{Parallel execution time}}$$

$$\text{Efficiency} = \frac{\text{Speedup}}{\text{Processors}}$$

# Amdahl's law


- Speedup is limited by the non-parallelizable serial portion of the work



<http://www.drdobbs.com/parallel/amdahls-law-vs-gustafson-barsis-law/240162980?pgno=2>



# Illustration of Amdahl Effect




# Review of Amdahl's Law

- Treats problem size as a constant
- Shows how execution time decreases as number of processors increases
- **Strong scaling**
  - Problem size is fixed
  - Number of processor increases

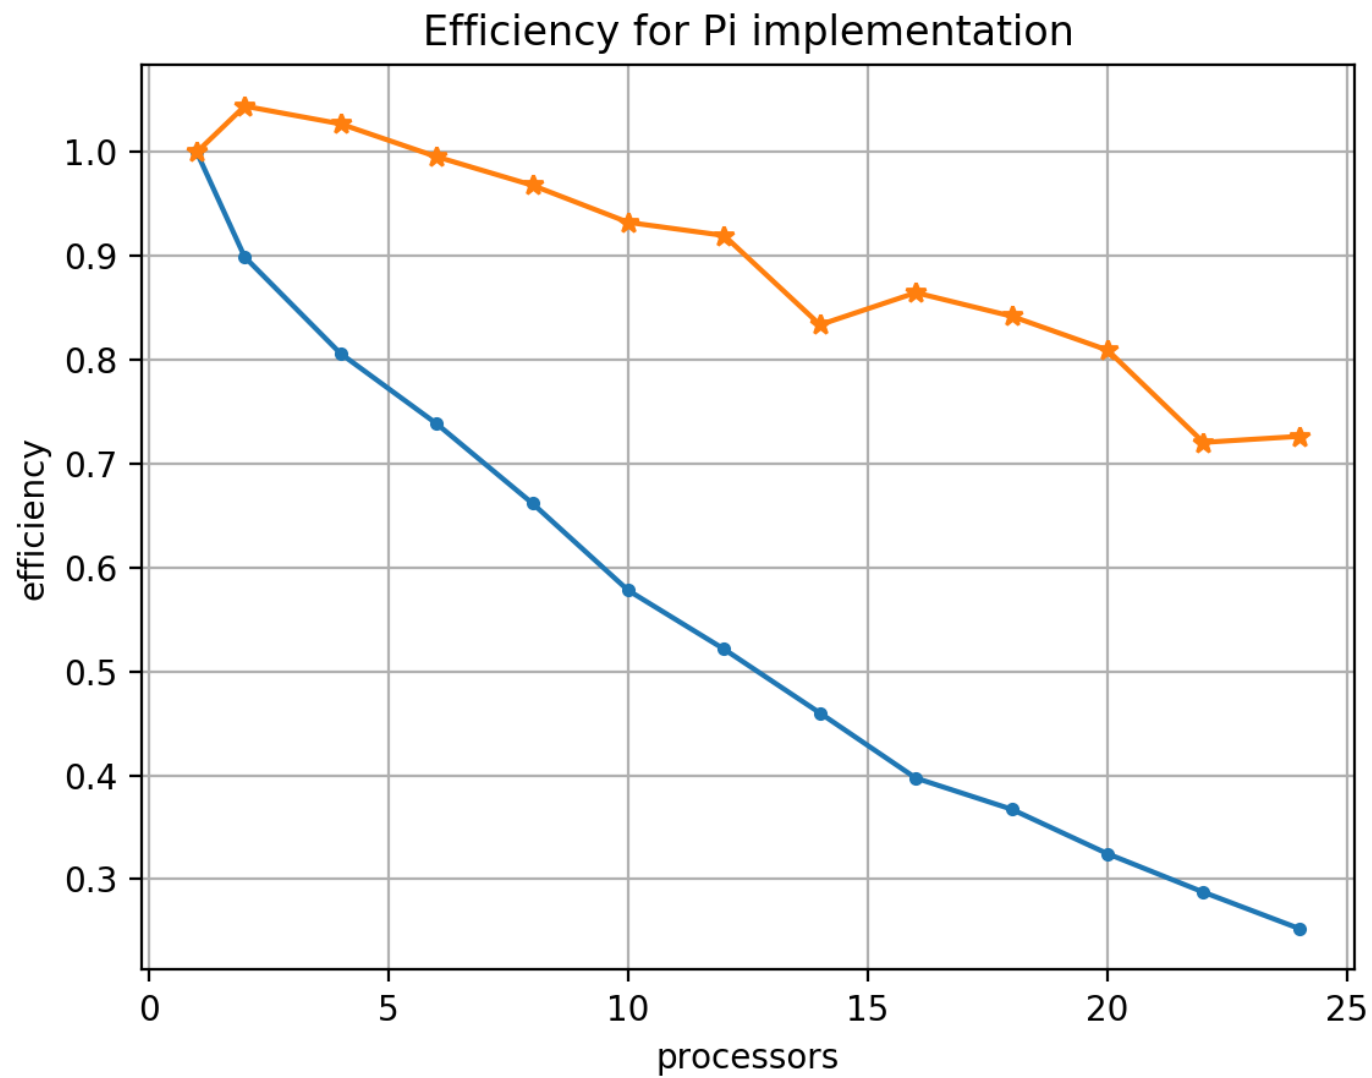
# Example

- Parallelizing Pi in python
  - Code in Day1/examples
  - pi\_serial.py
  - pi\_parallel.py
- Benchmarking on one node
  - 04-serial.sh
  - 04\_serial\_large.sh
- Strong Scaling
  - 04-strong\_scaling.sh
  - 04\_strong\_scaling\_large.sh
- Post processing
  - plot\_speedup.py

# Speedup




# Efficiency



# Another Perspective

- We often use faster computers to solve larger problem instances
- Let's treat time as a constant and allow problem size to increase with number of processors
- "...speedup should be measured by scaling the problem to the number of processors, not by fixing the problem size" – John Gustafson
- Weak scaling
  - Problem size per core stays constant
  - Overall program size increases with number of processors

# Gustafson-Barsis's Law



$$SS(N) = \frac{s+p*N}{s+p} = N + (1 - N)s$$

<http://www.drdoobbs.com/parallel/amdahls-law-vs-gustafson-barsis-law/240162980?pgno=2>

# Timing your code


- Use the unix time command
- Provides time for
  - Real – wall-clock time
  - User
  - Sys
- Use real time / wall-clock time for parallel benchmarking
- User + Sys gives you CPU time
  - Can be larger then real time – multithreading
- Use tools within your language
  - Module timeit for python



# Example Weak Scaling

- Run the following job scripts on Summit
  - 04-serial.sh
    - Creates timing for the serial time of the program
  - 04-weak\_scaling
    - Creates timing for weak scaling


# Weak Scaling



# Memory usage

- Massif as part of valgrind
- Valgrind help with memory related issues
  - `$ module load valgrind`
  - `$ valgrind --tool=massif python3 pi_serial.py 1000000`
  - `$ ms_print massif.out.88892`

# Massif Memory Usage Report



# Questions?

- Email [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Twitter: CUBoulderRC
- Link to survey on this topic:  
<http://tinyurl.com/curc-survey16>
- Slides:  
[https://github.com/ResearchComputing/Parallelization\\_Workshop/tree/master/Day1](https://github.com/ResearchComputing/Parallelization_Workshop/tree/master/Day1)