

Parallel Python Using Jupyter & IPyParallel

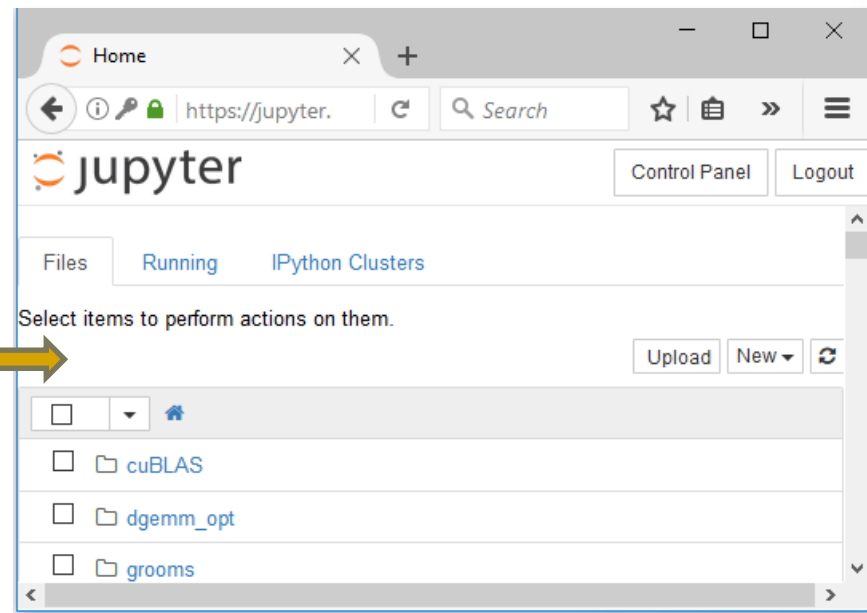
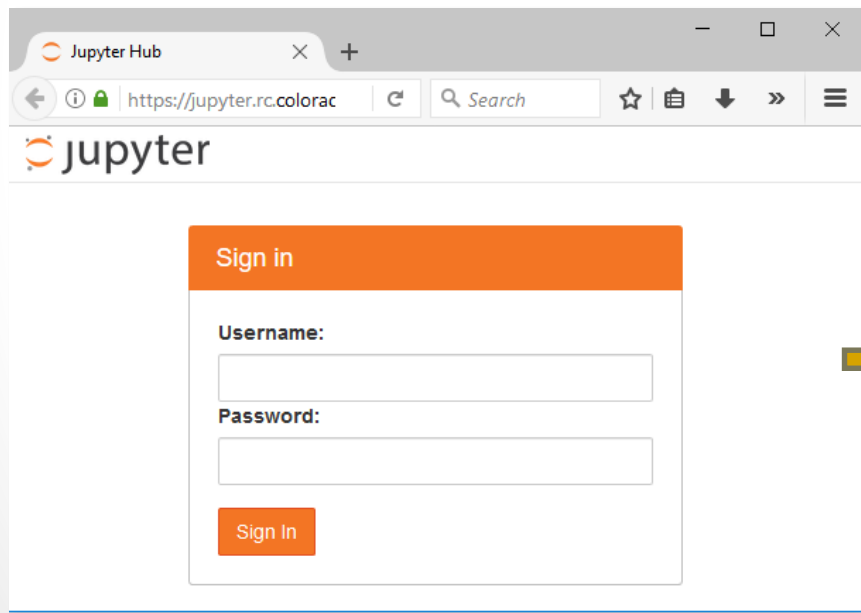
Nick Featherstone
CU Research Computing

[Web Link to These Slides](#)

Getting started...

- Login to the RC Jupyter Hub:

`https://jupyter.rc.colorado.edu`



Getting started...

Home × +

<https://jupyter.rc.colorado.edu/user/feathern/tree#ipyrc> Search

jupyter Control Panel Logout

Files Running **IPython Clusters**

IPython parallel computing clusters

profile	status	# of engines	action
default	stopped	<input type="text" value="1"/>	Start
crestone-cpu	stopped	4	Start
crestone-node	stopped	<input type="text" value="1"/>	Start
janus-cpu	stopped	<input type="text" value="1"/>	Start
janus-node	stopped	<input type="text" value="1"/>	Start

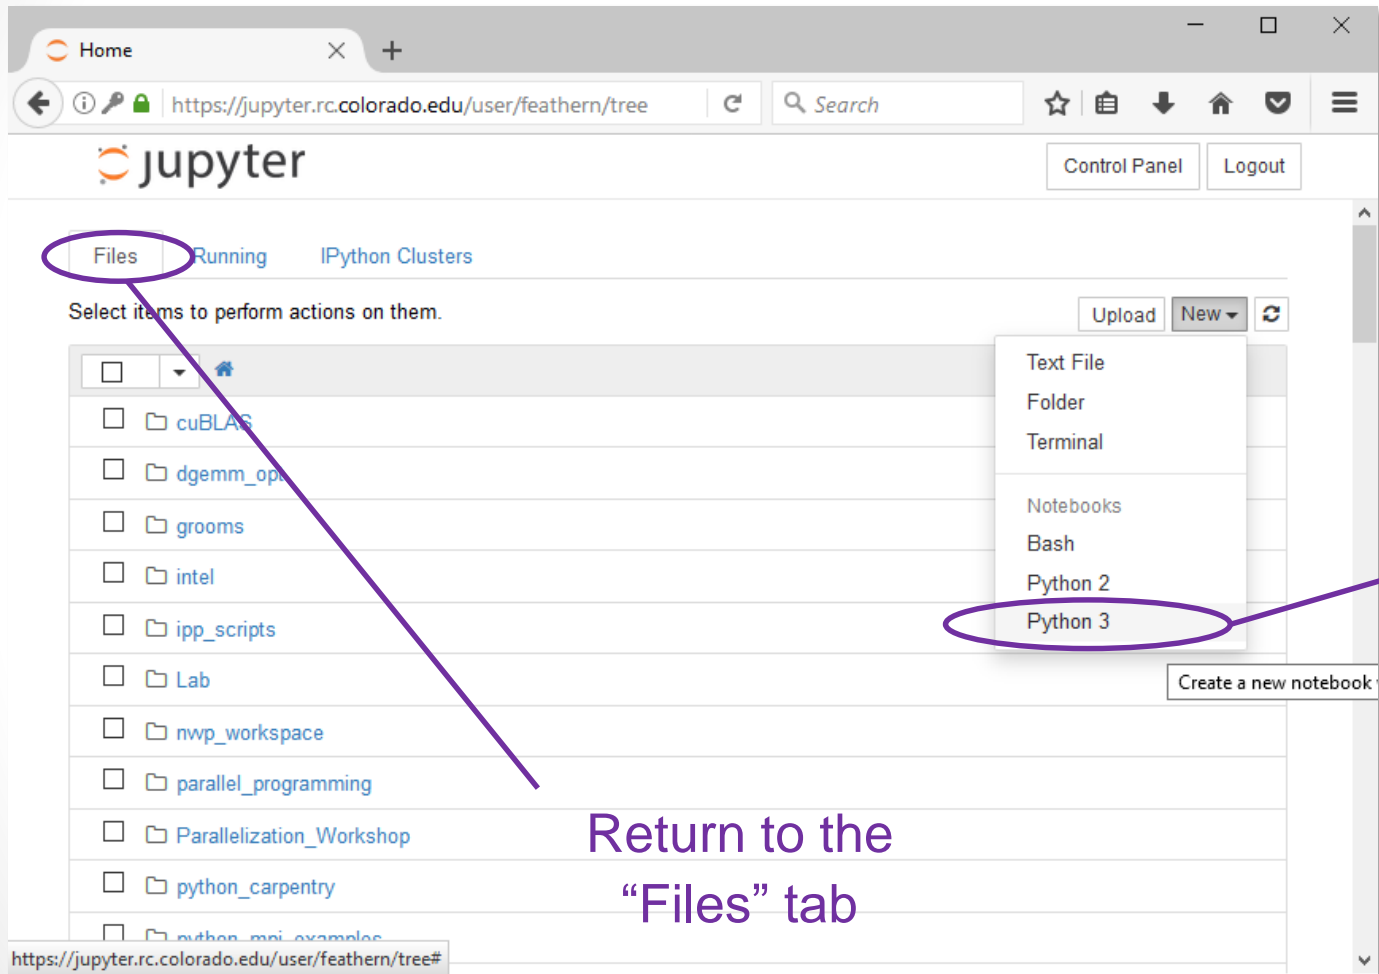
Select the "IPython Clusters" tab

Select "crestone-cpu"

4 engines

"start"

Getting started...




Start a Python 3 Notebook

Return to the
"Files" tab

Engine/Controller Paradigm

- What did we just do?
 - initiated a controller python session and 4 python engines
- **Most** of your code runs on the **controller (Jupyter notebook)**
- **Some** of your code runs on the **engines (the crestone cpu's)**



Documentation: <http://ipyparallel.readthedocs.io/en/latest/multiengine.html>

Hello World (1)

Most ipyparallel sessions will begin with something like:

- `import ipyparallel`
... enables access to ipyparallel features
- `rc=ipyparallel.Client(profile='crestone-cpu')`
... connects controller to the engines
- `print(rc.ids)`
... list of engine IDs, starting with 0

Open this file:

Parallelization Workshop / Day3-Parallel_Python /
session2_ipyparallel / examples /
hello.py

Hello World (2)

Some other interesting pieces that we will be working with:

```
all_proc = rc[:]
```

- ... “direct views” into each engine

- ... allows us to manipulate individual engines

```
hostnames = all_proc.apply_sync(socket.gethostname)
```

- ... tells all engines to call the `gethostname` function

- ... results stored in `hostnames`

Engine Views

The client object `rc` is essentially a list of engines.

We can explicitly reference individual engines or subsets of the engine pool by sampling or slicing `rc` as we would a list.

<code>engines = rc[:]</code>	view into all engines
<code>proc0 = rc[0]</code>	view into engine zero
<code>even = rc[range(0,4,2)]</code>	view into all even numbered engines

Variable Assignment

Create a variable 'a' on all engines:

```
engines = rc[:]  
engines['a'] = 2
```

Create a variable 'b' on a engine ID 1:

```
one = rc[1]  
one['b'] = 3
```

Initialize variable 'c' on ID 2 using 'b' from 1:

```
two = rc[2]  
two['c']=one['b']
```

[Open this file:](#)

Parallelization Workshop / Day3-Parallel_Python /
session2_ipyparallel / examples /
[assignment.py](#)

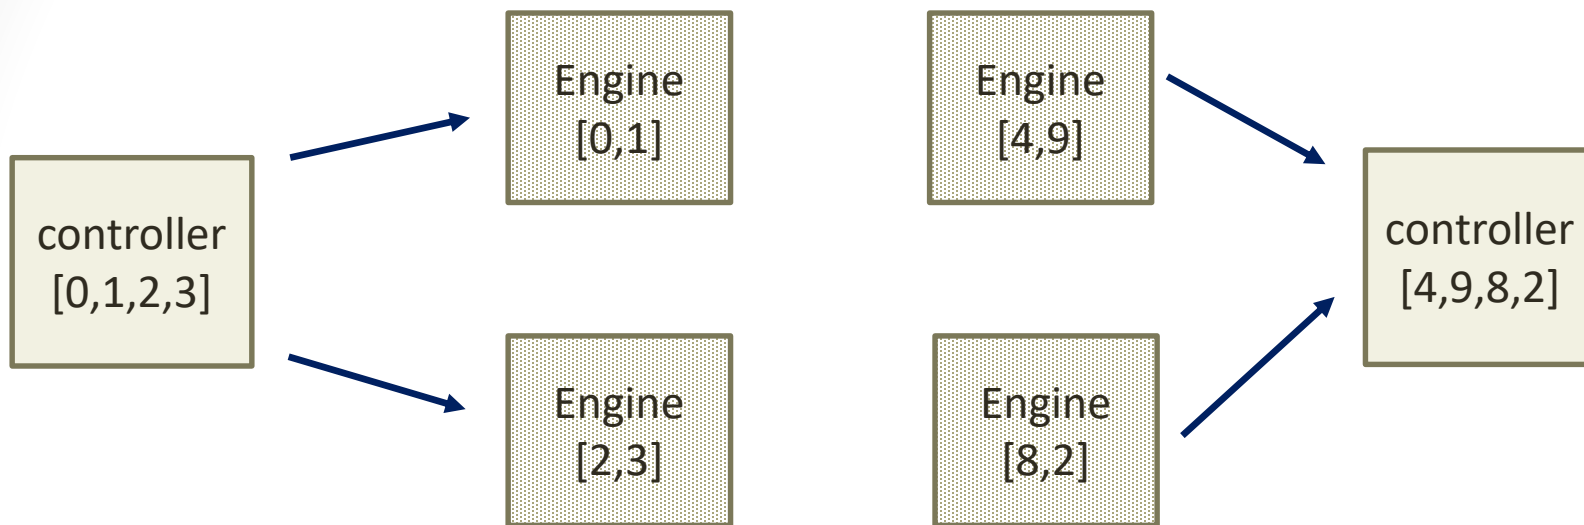
Exercise 1

Modify this program so that var1 receives the value n, where n is the remainder of the **Engine ID** (0 – 3) divided by 3.
i.e., $n = \text{Engine ID} \% 3$

Open this file:

Parallelization Workshop / Day3-Parallel_Python /
session2_ipyparallel / exercises /
assignment_ex.py

Scattering & Gathering



Scatter: distribute data from one process to the group

Gather: collate data from group onto one process

Both are common operations in parallel applications.

Scattering & Gathering

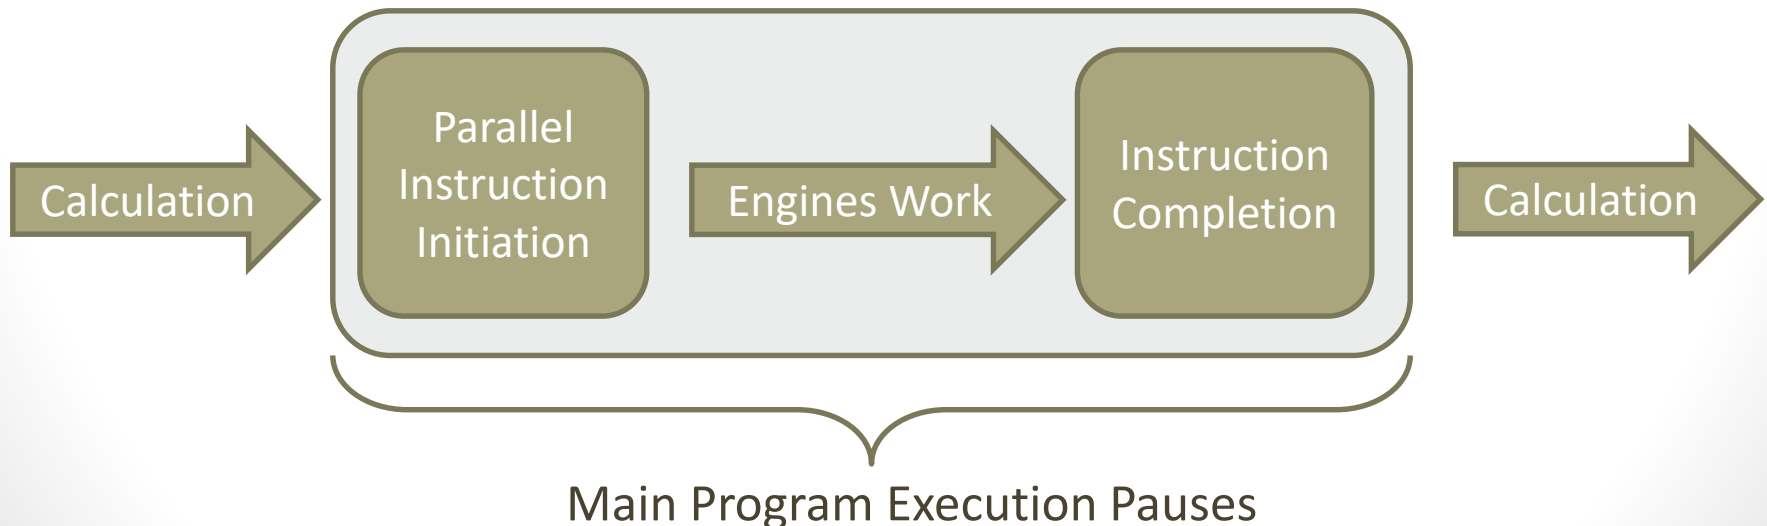
[Open this file:](#)

Parallelization Workshop / Day3-Parallel_Python /
session2_ipyparallel / examples /
[scatter_gather.py](#)

- To scatter list 'a' from the controller to the engines' variable 'mylist':
`all_proc.scatter('mylist' , a)`
- To view each engine's copy of 'mylist':
`sub_lists = all_proc['mylist']`
- To gather back from the engines to the controller:
`gathered = all_proc.gather('mylist')`

Quick Note: Blocking

- You may have noticed: `all_proc.block=True`
- This tells the engines and the controller to wait until parallel instructions have completed before resuming the code execution



Exercise 2

- Scatter list 'a' to all even processors, assigning its distributed values to the variable 'mylist' on each engine.
- Similarly, scatter list 'b' to all odd processors' 'mylist' variable
- Gather from all processors to create the list [0,1,0,1,2,3,4,9,....]


Open this file:

Parallelization Workshop / Day3-Parallel_Python /
session2_ipyparallel / exercises /

[scatter_gather_ex.py](#)

Who wants some Pi?

- Estimating pi is straightforward
- Compute n random coordinates in the domain: $-1 < (x,y) < 1$
- Count the number of coordinates m that fall within a unit circle centered on the origin
- Estimate = m/n



Function Evaluation via Map_Sync

- IPyParallel has an analog to map: `map_sync`

```
results= all_proc.map_sync(function_name, list_of_arguments)
```

- “results” contains a list of results from each process
- argument list is distributed (scattered) among all processes

Open this file:

Parallelization Workshop / Day3-Parallel_Python /
session2_ipyparallel / examples /

compute_pi.py

Exercises 3 & 4

- Try these exercises (instructions provided in each file):

Open this file:

Parallelization Workshop / Day3-Parallel_Python /
session2_ipyparallel / exercises /

parallel_functions.py
collatz.py

Classes in IPyParallel

- Class definitions exist solely on the controller/hub at first
- Definitions must be “pushed” to the engines before use:

```
all_proc.push(“engine class name” : controller class name)
```

Open this file:

Parallelization Workshop / Day3-Parallel_Python /
session2_ipyparallel / examples /
push.py