

Parallel Python Using MPI4Py


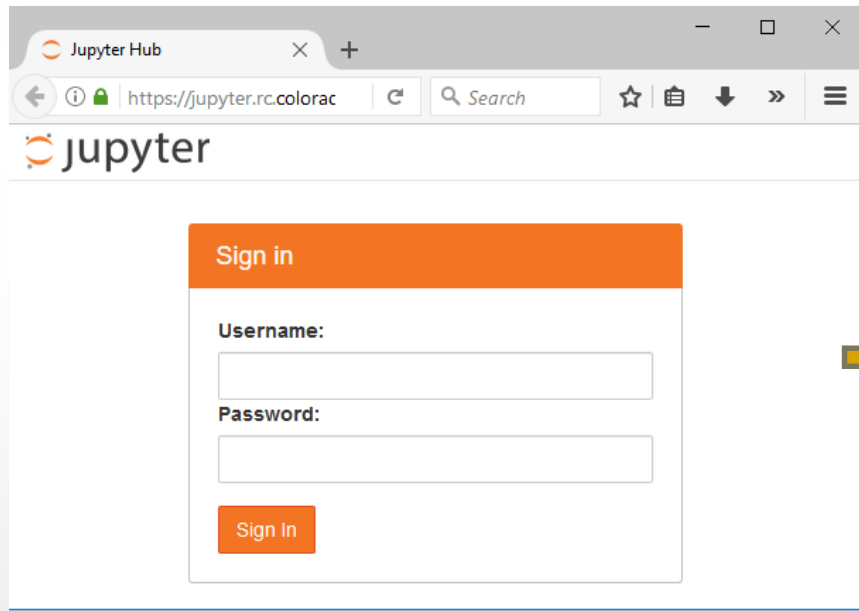
Nick Featherstone
CU Research Computing

[Web Link to These Slides](#)

Getting started...

Login to the RC Jupyter Hub:

`https://jupyter.rc.colorado.edu`



Getting started...

Home <https://jupyter.rc.colorado.edu/user/feathern/tree#ipyc> Search

jupyter Control Panel Logout

Files Running **IPython Clusters**

IPython parallel computing clusters

profile	status	# of engines	action
default	stopped	<input type="text" value="1"/>	Start
crestone-cpu	stopped	4	Start
crestone-node	stopped	<input type="text" value="1"/>	Start
janus-cpu	stopped	<input type="text" value="1"/>	Start
janus-node	stopped	<input type="text" value="1"/>	Start

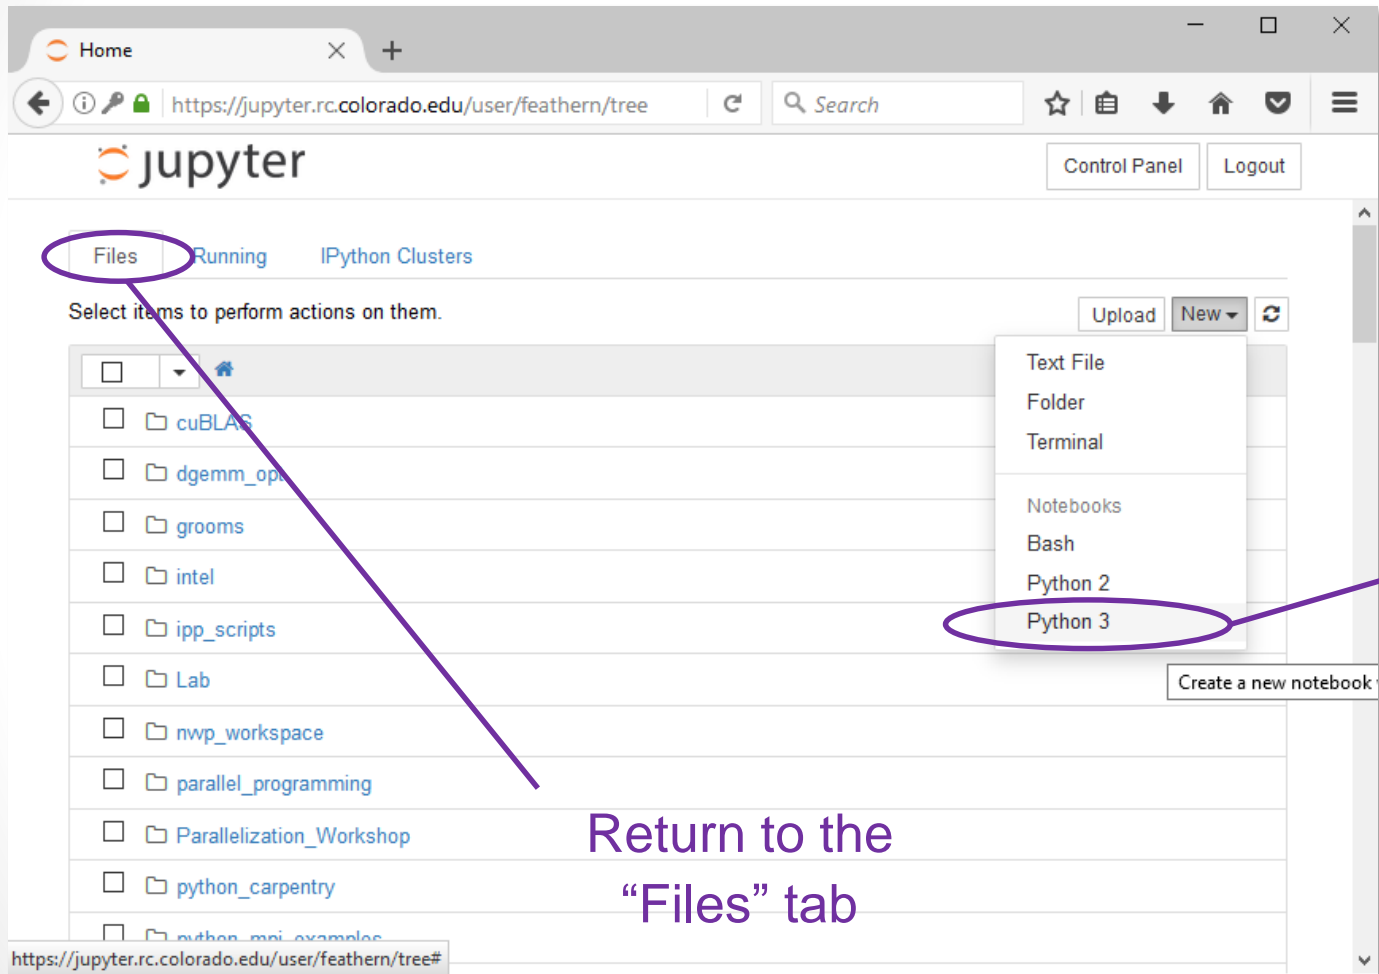
Select the "IPython Clusters" tab

Select "crestone-cpu"

4 engines

"start"

Getting started...




Start a Python 3 Notebook

Return to the
"Files" tab

Engine/Controller Paradigm


- What did we just do?
 - initiated a controller python session and 4 python engines
- **Most** of your code runs on the **controller (Jupyter notebook)**
- **Some** of your code runs on the **engines (the crestone cpu's)**



Documentation: <http://ipyparallel.readthedocs.io/en/latest/multiengine.html>

Message-Passing Paradigm

- No one is really in control.
- Everyone can communicate with each other
- **Entire code** runs on each process
- Processes referred to as MPI ranks



MPI4PY Documentation: <https://pythonhosted.org/mmpi4py/usrman/>

Message passing

- Most natural and efficient paradigm for distributed-memory systems
- Two-sided, **send** and **receive** communication between processes
- Efficiently portable to shared-memory or almost any other parallel architecture:
 - “assembly language of parallel computing” due to universality and detailed, low-level control of parallelism

More on message passing

- Provides natural synchronization among processes (through blocking receives, for example), so explicit synchronization of memory access is unnecessary
- Sometimes deemed tedious and low-level, but thinking about locality promotes
 - good performance,
 - scalability,
 - Portability
- Dominant paradigm for developing portable and scalable applications for massively parallel systems

Hello World

Open this file: **(DO NOT COPY/RUN THE PROGRAM YET)**

Parallelization Workshop / Day3-Parallel_Python /
session3_mpi4py / examples /
hello1.py

Every MPI program has two very important pieces

- Initialize communication

```
from mpi4py import MPI
```

- Finalize communication

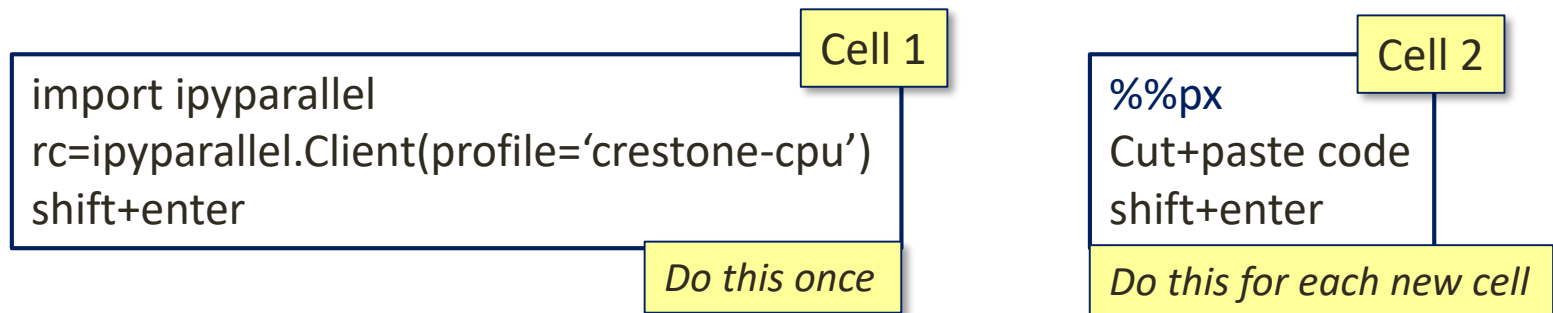
```
MPI.Finalize()
```

- The last very important part is commented out...

Logistics

In this workshop, we will be running our MPI code using the Jupyter notebook

- This is non-standard, but useful for an instructional setting
- We have to 'hack' things a bit to make ipyparallel and MPI work together
- Modified workflow:



- Normal codes invoking MPI Finalize are found in the *standard* directory
- DO NOT use those programs in today's session
- A sample batch script has been placed in the *standard* directory

Hello World

Most MPI programs have some other common components

- Define a communicator used to communicate within the process pool

```
comm = MPI.COMM_WORLD
```

- Identify number of processes

```
num_proc = MPI.COMM_WORLD.Get_size()
```

- Identify this process rank/ID

```
my_rank = MPI.COMM_WORLD.Get_rank()
```

- Recall that the entire code is running on each MPI rank
- num_proc will be the same on all ranks
- my_rank will receive a unique value for all ranks

[Open/Run this file:](#) Parallelization Workshop /
Day3-Parallel_Python /
session3_mpi4py / examples /
hello1.py

Hello World (2)

- Code is running on every process
- We don't always want every process performing I/O
- Common practice: use rank 0 in these situations
- Quick exercise:
 - Modify this program so that only even-numbered ranks print 'Hello from node...' message

```
if (my_rank == 0):  
    write stuff
```

[Open/Run this file:](#) Parallelization Workshop /
Day3-Parallel_Python /
session3_mpi4py / examples /
hello2.py

Flow Control using Barriers

- Different ranks may execute code at slightly different speeds
- This can lead to:
 - Jumbled output (everyone prints at once)
 - Race conditions (some ranks begin a task too soon)
- MPI's Barrier provides a method of synchronization

`comm.Barrier()`

- When a communicator's barrier method is called, all ranks in that communicator pause until everyone reaches that point in the code.


Flow Control using Barriers

[Open/Run this file:](#) Parallelization Workshop /
Day3-Parallel_Python /
session3_mpi4py / examples /
barrier.py

- Note the use of `sys.stdout.flush()`.
- If we don't flush the output buffer, the barrier will **appear** to be ineffective (try taking the flush out).
- EXERCISE:
Modify `exercises/barrier_ex.py` so that the different MPI ranks print 'hello from node...' in ascending order by rank.

Recall: Reduction

$A = [A[0], A[1], A[2], A[3]]$



- Sequential function evaluation using results from previous evaluations in tandem with new inputs
 - Think:
 - Adding numbers
 - Computing factorials
 - Common in parallel applications
- Scalar Result

Reduction in MPI

- Without the a dedicated controller, reduction requires some organized communication
- MPI's AllReduce functionality handles this 'under the hood'

Open/Run this file: Parallelization Workshop /
Day3-Parallel_Python /
session3_mpi4py / examples /
reduction.py

Reduction Syntax in MPI4PY

```
comm.Allreduce([local_sum, MPI.DOUBLE], [global_sum, MPI.DOUBLE], op=MPI.SUM)
```


- When performing an allreduce, we provide a few pieces of information
 - The communicator (implicitly stated via comm.Allreduce)
 - A list containing
 - The variable to be reduced across all processes
 - The TYPE of the reduction variable
 - A list containing
 - The variable in which to store the result
 - The TYPE of the result variable
 - The type of reduction we wish to perform
- For this workshop, we use two variable types:
 - MPI.DOUBLE : equivalent to 'float64' in numpy
 - MPI.INTEGER : equivalent to 'int32' in numpy
- Common reduction operations
 - MPI.MAX : computes the max
 - MPI.MIN : computes the min
 - MPI.SUM : computes the total

Exercise: Reduction

[Open/Run this file:](#) Parallelization Workshop /
Day3-Parallel_Python /
session3_mpi4py / exercises /
[collatz.py](#)


- Modify this code so that each rank uses Allreduce to compute the longest Collatz sequence occurring for numbers from 1 to 4000.

Reduction Application: Integrals



$$Area = \int_a^b f(x)dx \approx \sum_{i=1}^N A_i$$


Reduction Application: Integrals



$$Area = \int_a^b f(x)dx \approx \sum_{i=1}^N A_i$$

- Trapezoidal Rule

Load Balancing



- Idea: Assign each MPI rank a different range in x
- Total areas at the end via reduce

Exercise: Integration

Open/Run this file: Parallelization Workshop /
Day3-Parallel_Python /
session3_mpi4py / exercises /
trapezoid.py

Let have a look together...

- We've started the problem setup already:

```
//ntests = 10000; //uncomment  
ntrap = 1000000/num_proc;
```


- New local limits of integration are
myxone, myxtwo
- Use my_rank and num_proc to modify these limits appropriately...

```
xone = 1.0;  
xtwo = 2.0;  
// Each rank should integrate  
// What should deltax and myxo  
  
deltax = (xtwo-xone);  
myxone = xone;  
myxtwo = myxone+deltax;
```

... then run the code!

Message-Passing Example

[Open/Run this file:](#) Parallelization Workshop / Day3-
Parallel_Python /
session3_mpi4py / examples /
token_pass.py



Message Passing Questions

Which process is sending the message?

Where is the data on the sending process?

What kind of data is being sent?

How much data is there?

Which process is going to receive the message?

Where should the data be stored on the receiving process?

What amount of data is the receiving process prepared to accept?

Message Passing Syntax

