

Introduction to parallel computing with R

Thomas Hauser

Director of Research Computing

University of Colorado Boulder

https://github.com/ResearchComputing/Parallelization_Workshop

Creating parallel-ready R code

- Start with serial implementation
- Use function that are “parallel-friendly”
 - Use the apply family of functions
 - `base::apply`: Apply Functions Over Array Margins
 - `base::by`: Apply a Function to a Data Frame Split by Factors
 - `base::eapply`: Apply a Function Over Values in an Environment
 - `base::lapply`: Apply a Function over a List or Vector
 - `base::mapply`: Apply a Function to Multiple List or Vector Arguments
 - `base::rapply`: Recursively Apply a Function to a List
 - `base::tapply`: Apply a Function Over a Ragged Array
 - Easy to use parallel equivalents that will take care of a lot of things behind the scenes

Package parallel

- Based on multicore and snow
- Includes a parallel random number generator
- Supports single program multiple data paradigm (SPMD)
- Main interface is parallel version of lapply and similar
- Can use cores on a single machine (multicore) or several machines using MPI
- MPI support depends on the Rmpi package

Introduction to Parallel

- Scatter/gather paradigm
 - Scatter: breaking work into chunks and then distributes it to workers
 - Computation on a chunk
 - Gather: manager receives results and combines them to solution

Package: parallel

- Startup and stop
 - `makeCluster`: create a parallel cluster
 - `#` one or more parallel calls
 - `stopCluster`: stop a cluster
- Running a function on the cluster
 - `clusterCall`: calls a function `fun` with identical arguments ... on each node.
 - `mclapply(x, function, ..., mc.cores`
 - All objects and packages are automatically available
 - `parLapply(cl, x, function, ...)`
 - Have to explicitly export objects to the cluster
 - `clusterExport`
 - `clusterEvalQ`

Hello World – parallel version

- Run the following versions

```
$ sinteractive --partition=shas --qos=debug \  
--time=30:00 --ntasks=24 --nodes=1 \  
--reservation=parallelD4  
$ module purge  
$ module load R  
$ cd $HOME/Parallelization_Workshop/Day4-  
Parallel_R/examples/parallel  
$ Rscript hello_parallel.R
```

Sum the rows of a matrix

```
M <- matrix(1:, nrow = 3)
for (i in 1:3) {
  print(sum(M[i, ]))
}
```

- Order of loop is irrelevant
- This is a parallelizable algorithm

Functional programming in R

- Apply

```
M <- matrix(1:9, nrow = 3)
apply(M, 1, sum)
```

- [1] 12 15 18
- Functional programming in R
- Easier to parallelize
 - There are similar construct that work in parallel
- Independent parallel

Package parallel

- parApply: very similar to apply
- Simple approach to parallelization

Parallel version

```
library(parallel)
library(microbenchmark)
ncores = detectCores()
print(ncores)
cl <- makeCluster(ncores)
nrows <- 1000
M <- matrix(1:nrows^2, nrow = nrows, ncol =
nrows)
microbenchmark(parApply(M, 1, sum, cl=cl),
times=10)
stopCluster(cl)
```

Exercise: Parallel sum_rows


- Get a new node

```
$ Rscript sum_rows_apply_parallel.R
```


- Use a different number of cores
 - 4, 8, 12

Calculate PI with Monte Carlo

- Goal: estimate the area of a circle with radius = 1 and area = π using Monte Carlo integration.



Monte Carlo Integration



Monte Carlo with R

```
approx_pi <- function(n) {  
  # estimate pi w/ MC integration  
  x <- runif(n, min = -1, max = 1)  
  y <- runif(n, min = -1, max = 1)  
  V <- 4  
  f_hat <- ifelse(x^2 + y^2 <= 1, 1, 0)  
  V * sum(f_hat) / n  
}
```

apply() for vectors

sapply() returns vectors, matrices, and arrays

```
pi_hat <- sapply(n, approx_pi)  
str(pi_hat)
```

```
num [1:1000] 4 2.8 3.2 2.7 3.04 ...
```

Local parallelization

- Each MC integration is embarrassingly parallel!
- To parallelize:
 - start a cluster
 - compute simultaneously across the cluster
 - gather results
 - close cluster

the parallel package

```
cl <- makeCluster(2)

pi_hat <- parSapply(cl, n, approx_pi)

stopCluster(cl)
```

Parallel Random Numbers

- `RNGkind("L'Ecuyer-CMRG")`
- `set.seed(1.0)`
- Creates independent streams for each process

Exercise - pi_parallel.R

- Is the program using the weak scaling or strong scaling approach?
- Modify the program so that it uses the other approach.
- Run the program on your own node using
 - 4, 16 and 24 cores

Questions?

- Email rc-help@colorado.edu
- Twitter: CUBoulderRC
- Link to survey on this topic:
<http://tinyurl.com/curc-survey16>
- Slides:
https://github.com/ResearchComputing/Parallelization_Workshop

License

This work is licensed under the Creative Commons Attribution 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by/4.0/>

When attributing this work, please use the following text:
“OpenMP”, Research Computing, University of Colorado
Boulder, 2016. Available under a Creative Commons
Attribution 4.0 International License.

