

# Introduction to parallel computing with R - pbdMPI

Thomas Hauser

Director of Research Computing

University of Colorado Boulder


[https://github.com/ResearchComputing/Parallelization\\_Workshop](https://github.com/ResearchComputing/Parallelization_Workshop)

# Outline

- Distributed parallel computing
- Quick overview over MPI (Message Passing Interface)
- Package pbdMPI
- Examples
  - Hello World
  - Sum Rows
  - PI



# Distributed Memory Computer

- Processors have different content in memory
- Data exchange by message passing



# Programming Models

- Single Program  
Multiple Data  
(SPMD)
  - Same program runs on each process.
- Multiple Programs  
Multiple Data  
(MPMD)
  - Different programs runs on each process.



# Message passing

- Most natural and efficient paradigm for distributed-memory systems
- Two-sided, **send** and **receive** communication between processes
- Efficiently portable to shared-memory or almost any other parallel architecture:  
“assembly language of parallel computing” due to universality and detailed, low-level control of parallelism

# MPI standard

- MPI has been developed in three major stages
  - MPI 1 – 1994
  - MPI 2 – 1996
  - MPI 3 – 2012
- MPI Forum  
<http://www.mpi-forum.org/docs/docs.html>
- MPI Standard  
<http://www.mpi-forum.org/docs/mpi-3.0/mpi30-report.pdf>
- Using MPI and Using Advanced MPI  
<http://www.mcs.anl.gov/research/projects/mpi/usingmpi/>
- Online MPI tutorial  
<http://mpitutorial.com/beginner-mpi-tutorial/>

# MPI programs use SPMD model

- Same program runs on each process
- Build executable and link with MPI library
- User determines number of processes and on which processors they will run

# Execution

- You can run a MPI program with the following commands

```
$ mpirun -n 24 Rscript yourRprogram.R
```



# Run hello\_print.R

- Example – Run the hello\_print.R

```
$ sinteractive --partition=shas --qos=debug \  
--time=30:00 --ntasks=24 --nodes=1 \  
--reservation=parallelD4  
$ module purge  
$ module load R  
$ module load openmpi  
$ cd $HOME/Parallelization_Workshop/Day4-  
Parallel_R/examples/pbdMPI  
$ mpirun -n 10 Rscript hello_print.R
```
- Vary -n
- Is the output always in the same order?

# Programming in MPI

```
library(pbdMPI, quiet=TRUE) #include "mpi.h"
```

```
init()
```

```
.  
.   
.
```

```
finalize()
```

```
int ierr;
```

```
ierr = MPI_Init(&argc, &argv);
```

```
.  
.   
.
```

```
ierr = MPI_Finalize();
```

# MPI Communicator

- A collection of processors of an MPI program
- Used as a parameter for most MPI calls.
- Processors with in a communicator have a number
  - Rank: 0 to  $n-1$
- `comm`
  - Contains all processors of your program run
- You can create new communicators that are subsets
  - All even processors
  - The first processor
  - All but the first processor

# Programming in MPI

```
library(pbdMPI, quiet = TRUE)
```

```
init()
```

```
nprocs <- comm.size()
```

```
id <- comm.rank().
```

```
.
```

```
.
```

```
finalize()
```

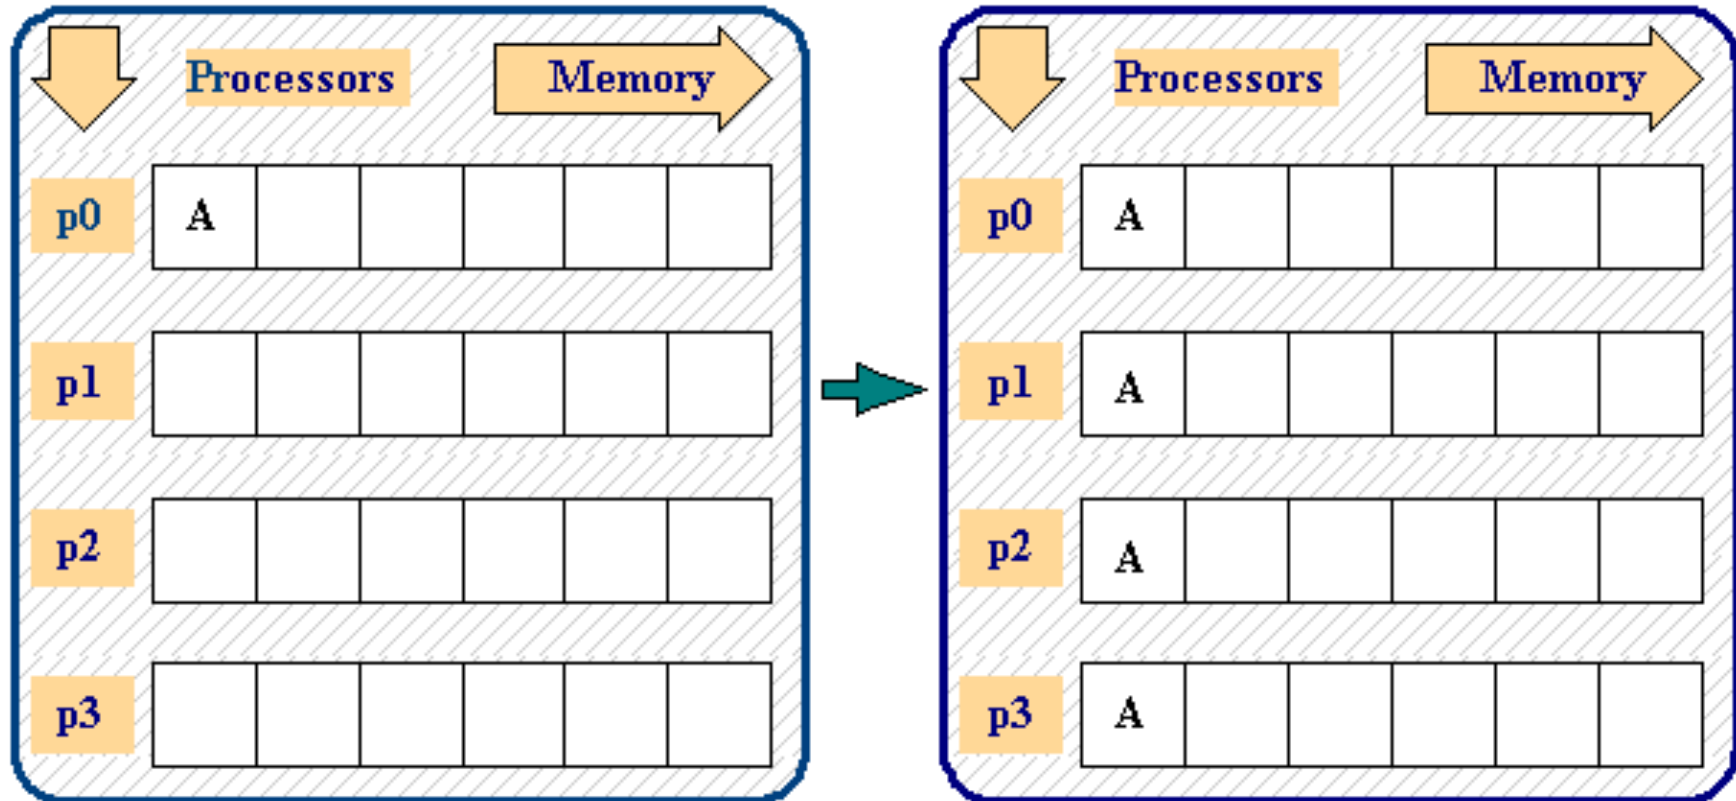
Determine process id or *rank* (here = id)

And number of processes (here = nprocs)

# Package pdbMPI

- Implements interface to MPI
  - > `comm.print(variable, all.rank=TRUE)`
  - > `comm.size()`
  - > `comm.rank()`
  - > `comm.set.seed(diff=TRUE)`
  - > `pbdApply(X, margin, func, ...)`
  - > `pbdLapply(X, func, ...)`
  - > `pbdSapply(X, fun, ...)`
  - > `bcast`
  - > `allgather`
  - > `reduce`


# Broadcast



```
send_count = 1;  
root = 0;  
MPI_Bcast ( &a, send_count, MPI_INT, root, comm )
```

Figure from MPI-tutor: <http://www.citutor.org/index.php>

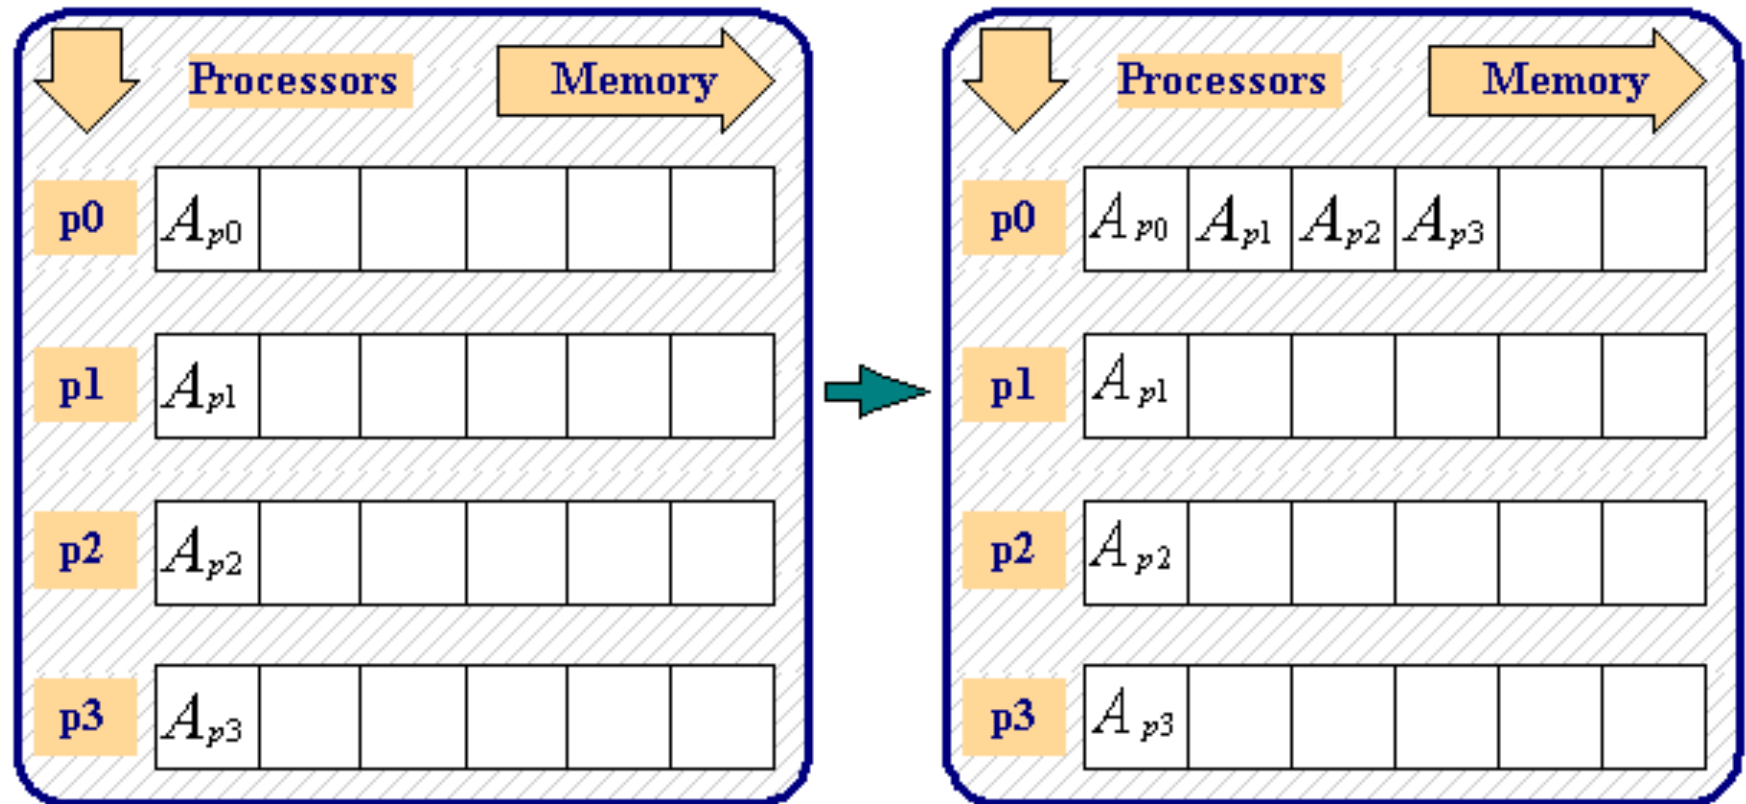
# Reduction



```
count = 1;  
rank = 0;  
MPI_Reduce ( &a, &x, count, MPI_REAL, MPI_SUM, rank, MPI_COMM_WORLD );
```

Figure from MPI-tutor: <http://www.citutor.org/index.php>


# Gather



```
send_count = 1;  
recv_count = 1;  
recv_rank = 0;  
MPI_Gather ( &a, send_count, MPI_REAL, &a, recv_count, MPI_REAL, recv_rank,  
MPI_COMM_WORLD );
```



# Scatter



```
recv_count = 1;  
send_rank = 0;  
MPI_Scatter ( &a, send_count, MPI_REAL,  
             &a, recv_count, MPI_REAL,  
             send_rank, MPI_COMM_WORLD );
```

# Printing

- `comm.print("String", all.rank=TRUE | FALSE)`
  - All processors have to participate
  - `all.rank=TRUE` – prints on all ranks
- Globally print or cat a variable from specified processors
- By default message is shown on screen
- **Warning: uses a barrier, so needs to be called by all processors**
  - **DEADLOCK danger**
  - Barrier is a synchronization between all processes. All processes have to join the call and processes wait until all have called it

# Deadlock

- Deadlock: process waiting for a condition that will never become true
- Easy to write send/receive code that deadlocks
  - Two processes: both receive before send
  - Send tag doesn't match receive tag
  - Process sends message to wrong destination process

# Run hello\_pbdMPI.R

- Example – Run the hello\_pbdMPI.R

```
$ module purge
$ module load R
$ module load openmpi
$ cd $HOME/Parallelization_Workshop/Day4-
Parallel_R/examples/pbdMPI
$ mpirun -n 10 Rscript hello_pbdMPI.R
```
- Vary -n
- Is the output always in the same order?

# Exercise – hello\_deadLock.R

- Run

```
$ mpirun -n 4 hello_deadLock.R
```
- What's happening?
- Try to fix the problem

# pbdApply, pbdSapply

```
$ pbdSapply(n, approx.pi, pbd.mode="spmd")
```

- `pbd.mode`
  - Single program multiple data – `spmd`
    - Need to distribute the data
      - Scatter or execute code on all processes
    - Need to collect the data
      - Gather or reduce
  - Master Worker – “`mw`”
    - Will distribute the first argument to the workers
    - Will get the result back
    - See later example

# Parallel Sum Rows

```
suppressMessages(library(pbdMPI , quietly = TRUE))
linit ()
.comm.size <- comm.size()
.comm.rank <- comm.rank()
comm.print(.comm.size,all.rank=TRUE)
comm.print(.comm.rank, all.rank=TRUE)
nrows <- 10
if (.comm.rank == 0) {
    M <- matrix(1:nrows^2, nrow = nrows, ncol = nrows)
}
r <- pbdApply(M, 1, sum, pbd.mode="mw", rank.source=0)
comm.print(r)
finalize
```


# Exercise: Sum Rows

- Run the program `sum_rows_apply_parallel.R` on 4 processes
- Modify the program
  - Replace “mw” with “spmd”
- What’s happening here



# Calculate PI with Monte Carlo

- Goal: estimate the area of a circle with radius = 1 and area =  $\pi$  using Monte Carlo integration.



# Exercise – pi\_pbdSapply

- Is the program using the weak scaling or strong scaling approach?
- Run the program on your own node using
  - 4, 16 and 24 cores
- Modify the program so that it uses the other approach.
- result variable has no value after the reduce
  - How can we fix this?

# Questions?

- Email [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Twitter: CUBoulderRC
- Link to survey on this topic:  
<http://tinyurl.com/curc-survey16>
- Slides:  
[https://github.com/ResearchComputing/Parallelization\\_Workshop](https://github.com/ResearchComputing/Parallelization_Workshop)

# License

This work is licensed under the Creative Commons Attribution 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by/4.0/>

When attributing this work, please use the following text:  
“OpenMP”, Research Computing, University of Colorado  
Boulder, 2016. Available under a Creative Commons  
Attribution 4.0 International License.

